

Introduction to Data Wrangling

Data Wrangling

- Real-world data is NOT clean!!
- We have to process it. More efficiency, better time use.
- This is where dplyr comes in: aka A Grammar of Data Manipulation.
- Hadley Wickham, 2014



DEE-PLIER



DIP-LER







DEEPLY-R



**DEE PEE
EL WHY ARE**





Note about installing dplyr

You can use the `install.packages` command to install the `dplyr` package.

After installing the package, you can load it into the workspace using the `library` command. Note that while you only need to install a package once, you need to load it into the workspace whenever you want to access it.

Alternatively, `dplyr` gets loaded automatically if you call `tidyverse` package, so you can skip this step. Tidyverse includes `ggplot2` too, and more.

```
> install.packages("dplyr")  
> library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following object is masked from 'package:MASS':
```

```
  select
```

```
The following objects are masked from 'package:stats':
```

```
  filter, lag
```

```
The following objects are masked from 'package:base':
```

```
  intersect, setdiff, setequal, union
```

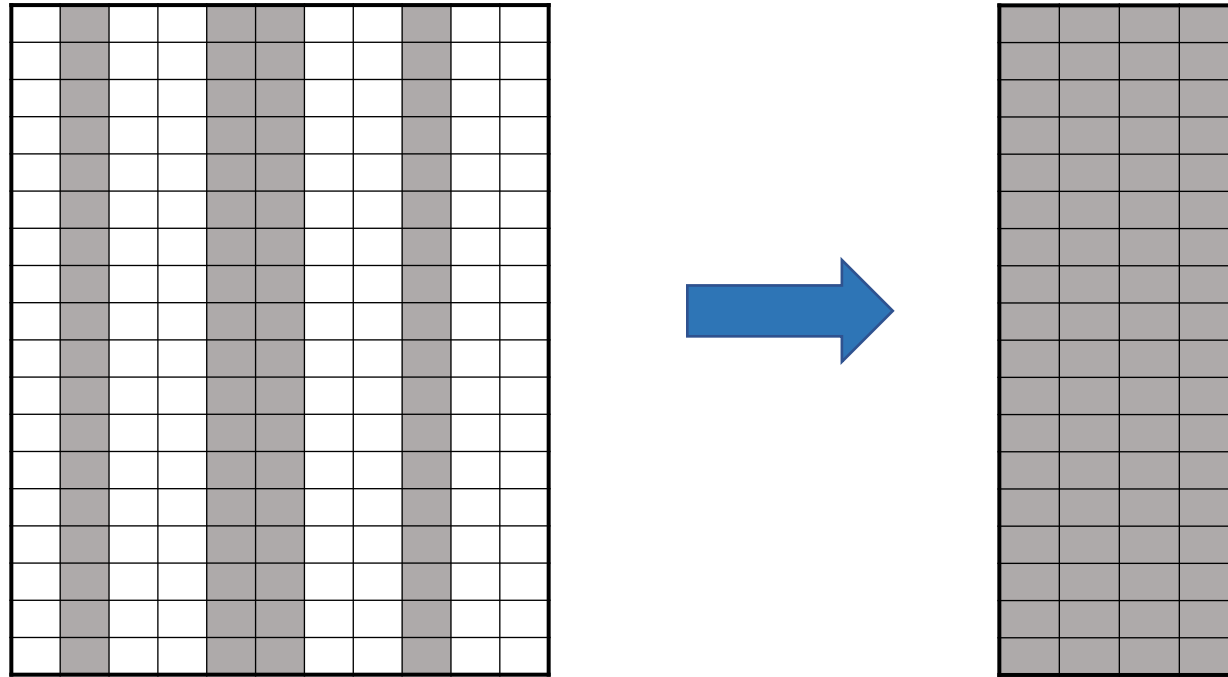
What is this telling us?

If we just type the object name, we get the object in the *last package loaded*. i.e., order matters. What is we want `select` from the `MASS` package?

```
> MASS::select
```

Our 1st function!: The `select` Function

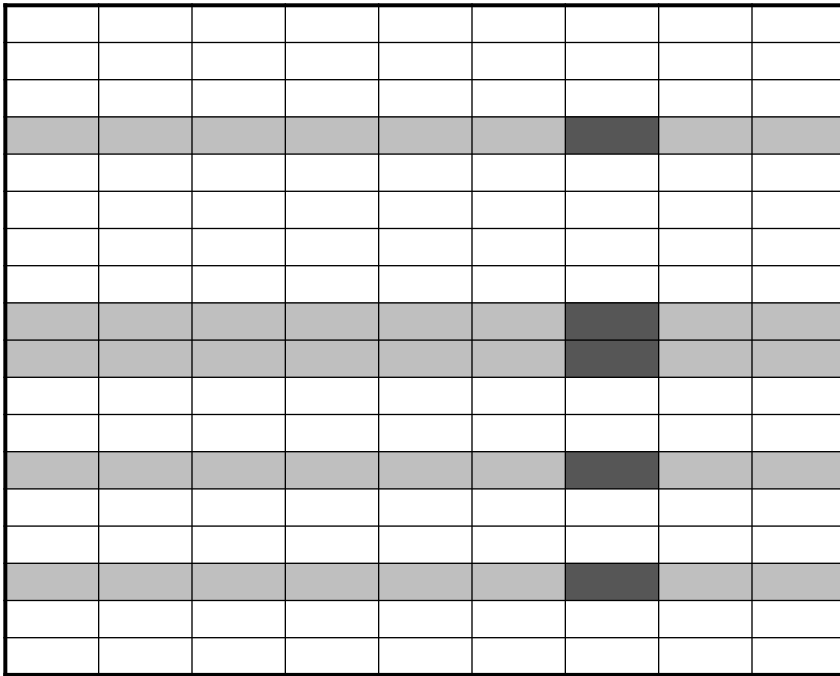
The `select` function allows you to create a new data frame that is a subset of an existing data frame by choosing a set of the columns of the original data frame.

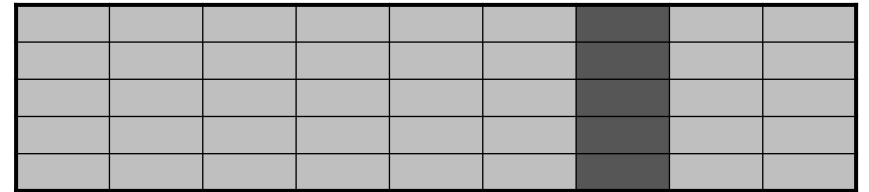


The general syntax is `select(data_frame, Var_name_1, Var_name_2, ...)`. Note that the entire column is transferred to the new data frame by the `select` function.

The `filter` Function

The `filter` function allows you to create a new data frame that is a subset of an existing data frame by choosing a set of the rows of the original data frame based on a collection of specified conditions.





The general syntax is `filter(data_frame, Condition,...)`. The condition can include logical operators for the row selection. Note that the entire row is transferred to the new data frame by the `filter` function.

Examples from the Cars93 Data Frame

Recall the `Cars93` data frame contains data for 93 cars sold in the U.S. during the year 1993. The first six rows of the data frame can be viewed using the `head` function.

```
> head(Cars93)
```

	Manufacturer	Model	Type	Min.Price	Price	Max.Price	MPG.city	MPG.highway	AirBags	DriveTrain
1	Acura	Integra	Small	12.9	15.9	18.8	25	31	None	Front
2	Acura	Legend	Midsize	29.2	33.9	38.7	18	25	Driver & Passenger	Front
3	Audi	90	Compact	25.9	29.1	32.3	20	26	Driver only	Front
4	Audi	100	Midsize	30.8	37.7	44.6	19	26	Driver & Passenger	Front
5	BMW	535i	Midsize	23.7	30.0	36.2	22	30	Driver only	Rear
6	Buick	Century	Midsize	14.2	15.7	17.3	22	31	Driver only	Front
	Cylinders	EngineSize	Horsepower	RPM	Rev.per.mile	Man.trans.avail	Fuel.tank.capacity	Passengers	Length	
1	4	1.8	140	6300	2890	Yes	13.2	5	177	
2	6	3.2	200	5500	2335	Yes	18.0	5	195	
3	6	2.8	172	5500	2280	Yes	16.9	5	180	
4	6	2.8	172	5500	2535	Yes	21.1	6	193	
5	4	3.5	208	5700	2545	Yes	21.1	4	186	
6	4	2.2	110	5200	2565	No	16.4	6	189	
	Wheelbase	Width	Turn.circle	Rear.seat.room	Luggage.room	Weight	Origin	Make		
1	102	68	37	26.5	11	2705	non-USA	Acura Integra		
2	115	71	38	30.0	15	3560	non-USA	Acura Legend		
3	102	67	37	28.0	14	3375	non-USA	Audi 90		
4	106	70	37	31.0	17	3405	non-USA	Audi 100		
5	109	69	39	27.0	13	3640	non-USA	BMW 535i		
6	105	69	41	28.0	16	2880	USA	Buick Century		

The select Function: Cars93 Example, and the %>% Operator

```
> Cars93_Ex1<-Cars93 %>% select(Type,EngineSize,DriveTrain,MPG.city)
> head(Cars93_Ex1)
```

	Type	EngineSize	DriveTrain	MPG.city
1	Small	1.8	Front	25
2	Midsize	3.2	Front	18
3	Compact	2.8	Front	20
4	Midsize	2.8	Front	19
5	Midsize	3.5	Rear	22
6	Midsize	2.2	Front	22

Notes:

- An arbitrary number of variable names can be passed into the function.
- The order of the variables in the new data frame matches the order they are entered into the function.
- The variable names are not in quotations in the function input.
- c() function is not used to combine column name objects.
- Keyboard shortcut for pipe: Ctrl + Shift + M (MacOS: Cmd + Shift + M)

The filter Function: Cars93 Example

```
> Cars93_Ex2<-Cars93 %>% filter(Type=="Small")
```

```
> head(Cars93_Ex2)
```

	Manufacturer	Model	Type	Min.Price	Price	Max.Price	MPG.city	MPG.highway	AirBags	DriveTrain	Cylinders
1	Acura	Integra	Small	12.9	15.9	18.8	25	31	None	Front	4
2	Dodge	Colt	Small	7.9	9.2	10.6	29	33	None	Front	4
3	Dodge	Shadow	Small	8.4	11.3	14.2	23	29	Driver only	Front	4
4	Eagle	Summit	Small	7.9	12.2	16.5	29	33	None	Front	4
5	Ford	Festiva	Small	6.9	7.4	7.9	31	33	None	Front	4
6	Ford	Escort	Small	8.4	10.1	11.9	23	30	None	Front	4
	EngineSize	Horsepower	RPM	Rev.per.mile	Man.trans.avail	Fuel.tank.capacity	Passengers	Length	Wheelbase		
1	1.8	140	6300	2890		Yes	13.2	5	177	102	
2	1.5	92	6000	3285		Yes	13.2	5	174	98	
3	2.2	93	4800	2595		Yes	14.0	5	172	97	
4	1.5	92	6000	2505		Yes	13.2	5	174	98	
5	1.3	63	5000	3150		Yes	10.0	4	141	90	
6	1.8	127	6500	2410		Yes	13.2	5	171	98	
	Width	Turn.circle	Rear.seat.room	Luggage.room	Weight	Origin	Make				
1	68	37	26.5	11	2705	non-USA	Acura Integra				
2	66	32	26.5	11	2270	USA	Dodge Colt				
3	67	38	26.5	13	2670	USA	Dodge Shadow				
4	66	36	26.5	11	2295	USA	Eagle Summit				
5	63	33	26.0	12	1845	USA	Ford Festiva				
6	67	36	28.0	12	2530	USA	Ford Escort				

The == is a test for equality and we need to enter the “Small” in quotations since we are filtering for rows where the Type variable is equal to “Small”.

Combining the `filter` and `select` Functions:

When we have several operations to complete that need to be nested, the *pipe operator* `%>%` can result in nice and readable code.

```
> Cars93_Ex8<-Cars93 %>%  
  filter(Type %in% c("Sporty","Compact") & Horsepower >=120) %>%  
  select(Model,Type,EngineSize,Cylinders,Horsepower,MPG.highway)  
> head(Cars93_Ex8)
```

	Model	Type	EngineSize	Cylinders	Horsepower	MPG.highway
1	90	Compact	2.8	6	172	26
2	Camaro	Sporty	3.4	6	160	28
3	Corvette	Sporty	5.7	8	300	25
4	LeBaron	Compact	3.0	4	141	28
5	Stealth	Sporty	3.0	6	300	24
6	Prelude	Sporty	2.3	4	160	31

The *pipe* above accomplishes the same set of operations as the code below (not so readable)

```
> Cars93_Ex7<-select(filter(Cars93, Type %in% c("Sporty","Compact") & Horsepower >=120),  
+ Model,Type,EngineSize,Cylinders,Horsepower,MPG.highway)
```

The %in% Operator

A long list of criteria can be cumbersome to type. We can use the %in% operator to check whether a value is in a list of possible values.

```
➤ Cars93_Ex7<-Cars93 %>%  
  filter(Type %in% c("Sporty","Compact") & Horsepower >= 120) %>%  
  select(Model,Type,EngineSize,Cylinders,Horsepower,MPG.highway)  
> head(Cars93_Ex7)
```

	Model	Type	EngineSize	Cylinders	Horsepower	MPG.highway
1	90	Compact	2.8	6	172	26
2	Camaro	Sporty	3.4	6	160	28
3	Corvette	Sporty	5.7	8	300	25
4	LeBaron	Compact	3.0	4	141	28
5	Stealth	Sporty	3.0	6	300	24
6	Prelude	Sporty	2.3	4	160	31

Here, the cars are returned that have a Type of either Sporty or Compact and have Horsepower greater than 120.

Practice time: msleep

Run `data(msleep)` to load `msleep` in, `?msleep` for the codebook

Use `filter()` to filter `msleep` to

1. only herbivores
2. any animal that is awake for at least 12 hours a day
3. only herbivores + that are awake for at least 12 hours of a day.
4. Name and awake columns for only herbivores + that are awake for at least 12 hours of a day.

Do it using base R, and `dplyr`.

Hint: if you get a slightly different output, try adding `%>% drop_na(colname)`

Practice `%in%`:

Return herbivores and carnivores that sleep at least 12 hours a day

Do it using `%in%` and `"|"` which is the symbol for “or”

**everything you do in tidyverse, I
can do in data.table**



**everything you do in data.table,
I can do in base R**



**everything you do in base R,
I can do with pen and paper**

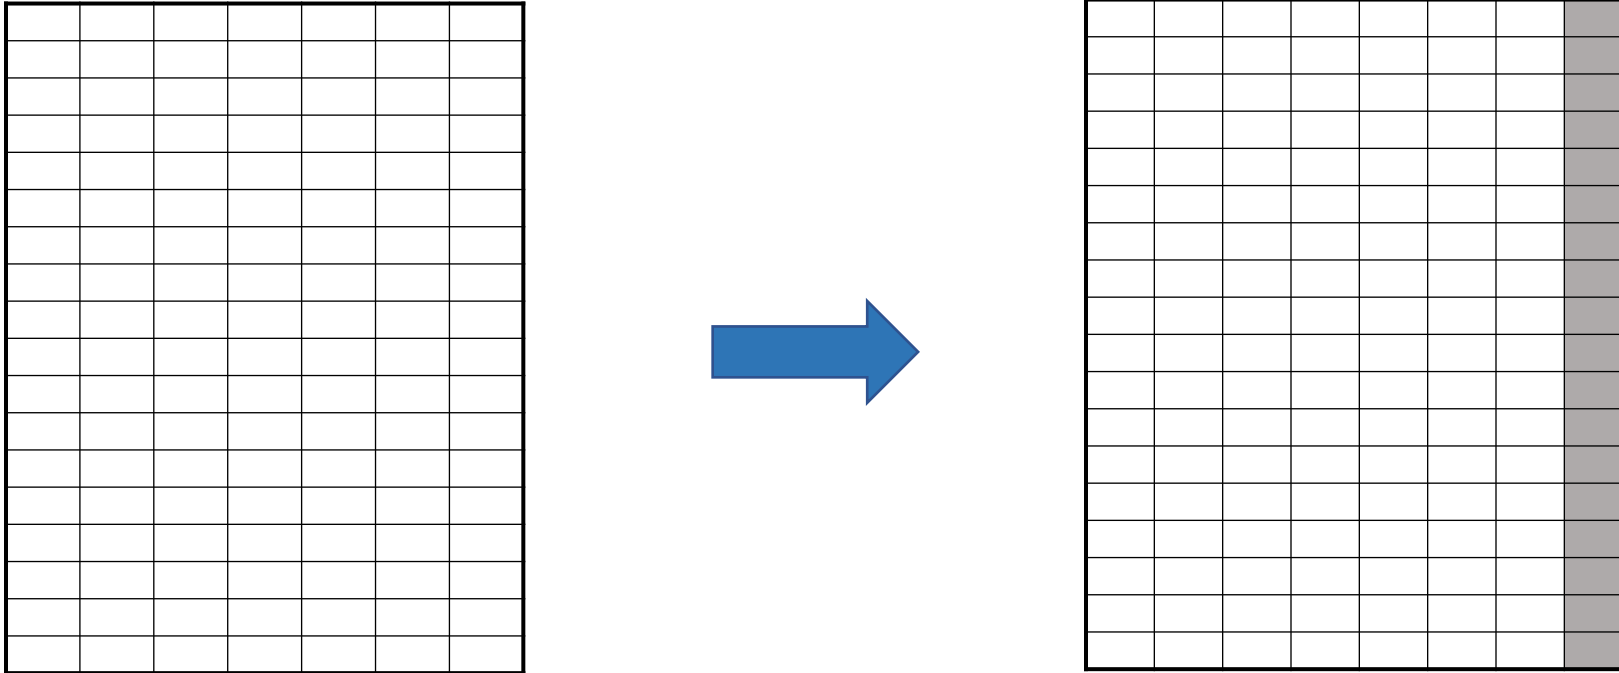


**arghbrbr@#&zzzegh%%~!!!!
I am the best!!!**



The mutate Function

The `mutate` function allows you to create a new data frame consisting of the original data frame with a column appended on the right end.



The general syntax is `mutate(data_frame, Var_Name = function(...))`. The new column will have the name given in `Var_Name` and be computed from the function provided.

Example of the mutate Function

We now add a column to the `Cars93` data frame containing the horsepower per liter of engine size. That is, we define the new variable `HPpLiter` by the function

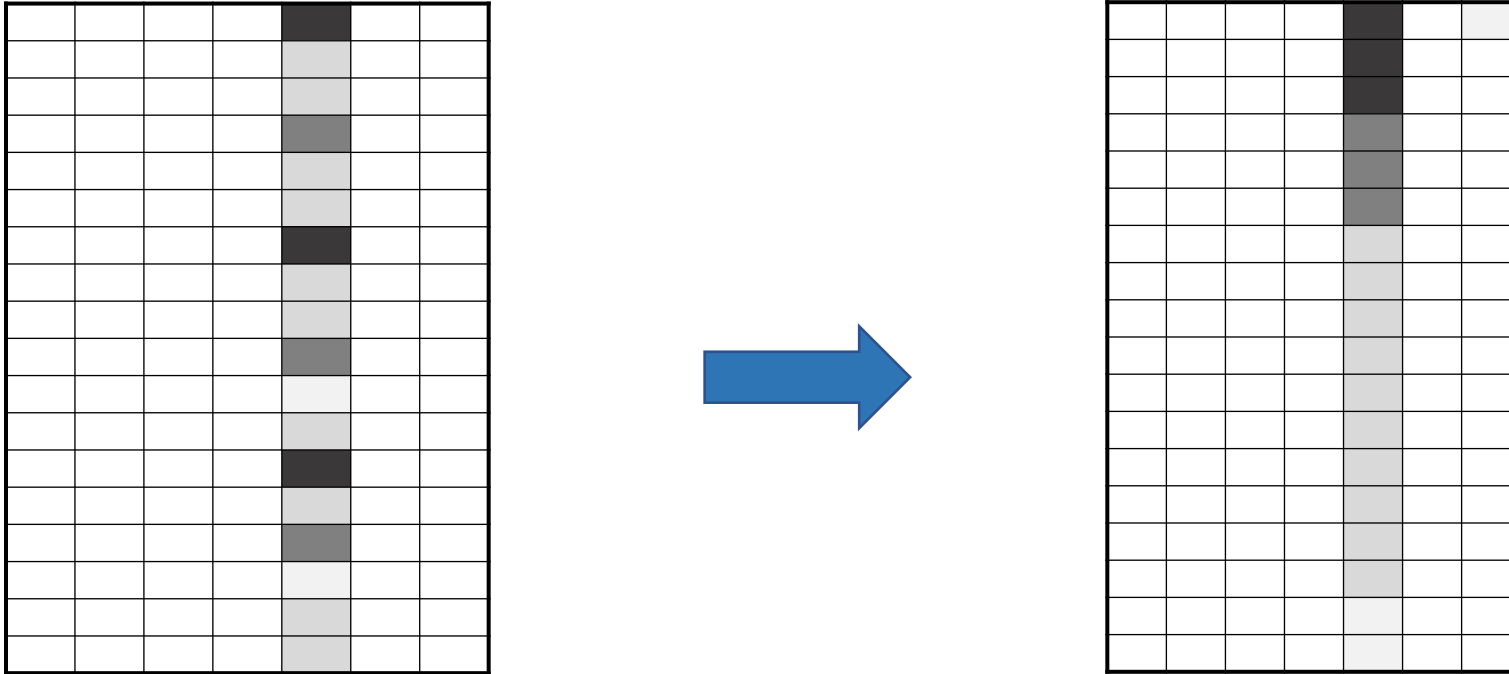
$$\text{HPpLiter} = \frac{\text{Horsepower}}{\text{EngineSize}}$$

```
> Cars93_Ex9<-Cars93 %>%  
  mutate(HPpLiter=Horsepower/EngineSize) %>%  
  select(Model,Type,EngineSize,Cylinders,Horsepower,MPG.highway,HPpLiter)  
> head(Cars93_Ex9)
```

	Model	Type	EngineSize	Cylinders	Horsepower	MPG.highway	HPpLiter
1	Integra	Small	1.8	4	140	31	77.77778
2	Legend	Midsize	3.2	6	200	25	62.50000
3	90	Compact	2.8	6	172	26	61.42857
4	100	Midsize	2.8	6	172	26	61.42857
5	535i	Midsize	3.5	4	208	30	59.42857
6	Century	Midsize	2.2	4	110	31	50.00000

The `arrange` Function

The `arrange` function allows you to sort a data frame by variable producing a new data frame ordered by that variable.



The general syntax is `arrange(data_frame, Var_Name1, Var_Name2, ...)`. The rows are sorted by the variables selected in the order they are entered. That is, on `Var_Name1` first, with ties broken by `Var_Name2`, etc. The default is ascending order, and descending order can be selected by using `desc(Var_Name1)` when passing in the variable name.

Example of the arrange Function

In the previous example we used the `mutate` function to add a column to the `Cars93` data frame containing the horsepower per liter. We now sort this new data frame first by the number of cylinders and then by the engine's horsepower.

```
> Cars93_Ex10<-Cars93 %>%  
  mutate (HPpLiter=Horsepower/EngineSize) %>%  
  select (Model,Type,EngineSize,Cylinders,Horsepower,MPG.highway,HPpLiter) %>%  
  arrange (Cylinders,desc (Horsepower) )  
> head(Cars93_Ex10)
```

	Model	Type	EngineSize	Cylinders	Horsepower	MPG.highway	HPpLiter
1	Justy	Small	1.2	3	73	37	60.83333
2	Swift	Small	1.3	3	70	43	53.84615
3	Metro	Small	1.0	3	55	50	55.00000
4	535i	Midsize	3.5	4	208	30	59.42857
5	626	Compact	2.5	4	164	34	65.60000
6	Prelude	Sporty	2.3	4	160	31	69.56522

Observe that we used `desc (Horsepower)` to sort on this variable from highest to lowest.

The rename Function

The `rename` function allows you to rename columns in a data frame.

The syntax is `rename(data_frame, New_Name1=Old_Name1, New_Name2=Old_Name2 ...)`. The variable named `Old_Name` is renamed with the value entered for `New_Name` in the output data frame. Rows are not affected, and the column order remains unchanged.

Example: We can rename two of the columns in the data frame from the previous example using the `rename` function as follows:

```
> Cars93_Ex11<- Cars93_Ex10 %>%  
  rename(HP=Horsepower, MPG_Highway=MPG.highway)  
> head(Cars93_Ex11)
```

	Model	Type	EngineSize	Cylinders	HP	MPG_Highway	HPpLiter
1	Justy	Small	1.2	3	73	37	60.83333
2	Swift	Small	1.3	3	70	43	53.84615
3	Metro	Small	1.0	3	55	50	55.00000
4	535i	Midsize	3.5	4	208	30	59.42857
5	626	Compact	2.5	4	164	34	65.60000
6	Prelude	Sporty	2.3	4	160	31	69.56522

More Examples with the `mutate` Function

Suppose we are interested in the engines of the various cars. We previously used the `mutate` function to add a column containing the horsepower per liter of engine size.

```
> head(Cars93_Ex11)
```

	Model	Type	EngineSize	Cylinders	HP	MPG_Highway	HPpLiter
1	Justy	Small	1.2	3	73	37	60.83333
2	Swift	Small	1.3	3	70	43	53.84615
3	Metro	Small	1.0	3	55	50	55.00000
4	535i	Midsize	3.5	4	208	30	59.42857
5	626	Compact	2.5	4	164	34	65.60000
6	Prelude	Sporty	2.3	4	160	31	69.56522

What if we wanted the data in the `HPpLiter` column to be rounded to the nearest tenth? (Note that we could have done this when we first added the column, but we forgot). We can use the `mutate` function to accomplish this as well.

```
➤ Cars93_Ex12<- Cars93_Ex11 %>%  
  mutate(HPpLiter=round(HPpLiter,1))
```

Observe that we have entered a variable name that is already in the data frame. How does the `mutate` function process this request?

More Examples with the mutate Function

Viewing the function output produces

```
> head(Cars93_Ex12)
```

	Model	Type	EngineSize	Cylinders	HP	MPG_Highway	HPpLiter
1	Justy	Small	1.2	3	73	37	60.8
2	Swift	Small	1.3	3	70	43	53.8
3	Metro	Small	1.0	3	55	50	55.0
4	535i	Midsize	3.5	4	208	30	59.4
5	626	Compact	2.5	4	164	34	65.6
6	Prelude	Sporty	2.3	4	160	31	69.6

The `mutate` function did not append a new column on the end of the data frame, but rather **replaced** the existing column with the matching name with the updated information.

Suppose that engines producing 65 horsepower or more per liter of displacement can be considered high-performance while those producing less than this value are regular performance. How can we add a column indicating this to our data frame?

More Examples with the mutate Function

We can combine the `mutate` function with the `ifelse` function to obtain this result

```
> Cars93_Ex13<- Cars93_Ex12 %>%  
  mutate(Performance=(HPpLiter>=65)*1)  
> head(Cars93_Ex13)
```

	Model	Type	EngineSize	Cylinders	HP	MPG_Highway	HPpLiter	Performance
1	Justy	Small	1.2	3	73	37	60.8	0
2	Swift	Small	1.3	3	70	43	53.8	0
3	Metro	Small	1.0	3	55	50	55.0	0
4	535i	Midsize	3.5	4	208	30	59.4	0
5	626	Compact	2.5	4	164	34	65.6	1
6	Prelude	Sporty	2.3	4	160	31	69.6	1

Note that the condition inside returns a boolean vector, which can be multiplied by a number to get integer output. The `mutate` function appends the result on the right side of the data frame. We can obtain further information using

```
> table(Cars93_Ex13$Performance)
```

```
0  1  
79 14
```


The summarize Function

- The `summarize` function allows you to produce a data frame with user chosen statistics calculated from the columns of the input data frame.
- Often used with the `group_by` function which allows the statistics to be computed for particular groups in the input data frame.

Example: We have the data on the engines in the `Cars93_Ex14` data frame.

```
> head(Cars93_Ex14)
```

	Model	Type	EngineSize	Cylinders	HP	MPG_Highway	HPpLiter	Performance
1	Justy	Small	1.2	3	73	37	60.8	Regular
2	Swift	Small	1.3	3	70	43	53.8	Low
3	Metro	Small	1.0	3	55	50	55.0	Regular
4	535i	Midsize	3.5	4	208	30	59.4	Regular
5	626	Compact	2.5	4	164	34	65.6	High
6	Prelude	Sporty	2.3	4	160	31	69.6	High

We can use the `summarize` function to obtain information about the engines grouped by, for example, the car `Type` or the `Cylinders` variable.

The summarize Function

An example of the `summarize` function is provided in the following pipeline:

```
> Engine_Summary<-Cars93_Ex14 %>%  
  group_by(Type) %>%  
  summarize(Num=n(),Min_Size=min(EngineSize),Max_Size=max(EngineSize),Ave_HP=mean(HP),  
    Median_MPG_HWY=median(MPG_Highway),Num_High_Performance=sum(Performance=="High"))  
  `summarize()` ungrouping output (override with `.groups` argument)  
> Engine_Summary
```

	Type	Num	Min_Size	Max_Size	Ave_HP	Median_MPG_HWY	Num_High_Performance
1	Compact	16	2.0	3.0	131.0000	30.0	4
2	Large	11	3.3	5.7	179.4545	26.0	0
3	Midsize	22	2.0	4.6	173.0909	26.5	4
4	Small	21	1.0	2.2	91.0000	33.0	5
5	Sporty	14	1.3	5.7	160.1429	28.5	3
6	Van	9	2.4	4.3	149.4444	22.0	0

The code above creates five new variables of interest and calculates them for the groups determined by the `group_by` function. The `n()` function gives the number of units in the group. The output of the function is a tibble we convert to a data frame using the `data.frame` function.

Example of the summarize Function

A second example of the `summarize` function is provided in the following pipeline:

```
> Engine_Summary_2<-Cars93_Ex14 %>%  
group_by(Cylinders) %>%  
summarise(Num=n(),Min_Size=min(EngineSize),Max_Size=max(EngineSize),Ave_HP=mean(HP),  
Median_MPG_HWY=median(MPG_Highway),Num_High_Performance=sum(Performance=="High"))  
`summarise()` ungrouping output (override with `.groups` argument)  
> Engine_Summary_2
```

	Cylinders	Num	Min_Size	Max_Size	Ave_HP	Median_MPG_HWY	Num_High_Performance
1	3	3	1.0	1.3	66.0000	43.0	0
2	4	49	1.3	3.5	113.4694	31.0	10
3	5	2	2.4	2.5	138.5000	24.5	1
4	6	31	2.8	5.7	175.5806	26.0	4
5	8	7	4.5	5.7	234.7143	25.0	0
6	rotary	1	1.3	1.3	255.0000	25.0	1

The same five new variables of interest are calculated but this time the engines are grouped by the `Cylinders` variable. Again, we which we convert the tibble output to a data frame using the `data.frame` function.

Tibbles

A core component of the tidyverse is the tibble. Tibbles are a modern rework of the standard `data.frame`, with some internal improvements to make code more reliable. They are similar to data frames, but do not follow all of the same rules. For example, tibbles can have numbers/symbols for column names, which is not allowed in base R.

If you use a function from the `dplyr` package that returns a tibble as output, you can convert it to a data frame using the `data.frame` function. For example

```
> Result
```

```
# A tibble: 3 x 4
```

	cyl	N	Mean_HP	Mean_mpg
	<dbl>	<int>	<dbl>	<dbl>
1	4	11	82.6	26.7
2	6	7	122.	19.7
3	8	14	209.	15.1

```
> Result<-data.frame(Result)
```

```
> Result
```

	cyl	N	Mean_HP	Mean_mpg
1	4	11	82.63636	26.66364
2	6	7	122.28571	19.74286
3	8	14	209.21429	15.10000

Practice time: msleep

Return average awake time by `vore` sorted by average awake time in descending order.

Create a new column `brain_per` that shows the percentage of body weight that brain takes (use the ratio `brainwt/bodywt`)

What is the average `brain_per`? (Hint: use `na.rm=T`) What does its distribution look like? (Hint: you can combine `dplyr` with `ggplot2` in one line!)

Return the top 5 animals with the highest `brain_per`.

Make a scatterplot showing `brain_per` against `awake`.

Make a scatterplot showing `brain_per` against `awake` with different colors by `vore`.